

Operating Systems – Project 1

- Simple Scheduling -

ChangYoon Lee, ChanHo Park

Dankook University, Korea Republic
32183641@gmail.com, chanho.park@dankook.ac.kr

Abstract. Most modern operating systems have extended the concept of the process to allow having multiple threads of execution and thus to perform more than one task at a time (*Silberschatz et al., 2014*). While allowing this multi-threaded concept, there are some issues to consider in designing it: Synchronization (Concurrency Control). In this paper, we will discuss the concept of the multi-thread, synchronization examples, and their solutions, especially on producer and consumer problems. Also, we will implement the example case of producer and consumer problems with a multi-threaded word count program and evaluate it.

Keywords: Concurrency Control, Multi-Thread, Mutex, Producer and Consumer, Semaphore, Synchronization, Reader and Write

1 Introduction

In a single-processor computer system, only one process can run at a single time. Other processes must wait until the CPU resources are free and can be rescheduled. A process is executed until it must wait, typically for the completion of the I/O request. However, in multiprogramming, some process runs at all time, to maximize CPU utilization. Multiprogramming tries to use the waiting time productively. Some process is loaded into the memory at one time, and when one process has to wait, the operating system takes the CPU resources away from the process and gives the CPU resources to another process. The following progress continues, every time one process has to wait, while another process takes over the use of the CPU resources.

Scheduling the following progress is a fundamental operating system function. Almost all the computer resources are scheduled before use. Since the CPU is one of the primary computer resources, CPU scheduling is central to the operating system design.

In this paper, we will first explain the concepts of process, process scheduling, Inter-process communication (IPC), and CPU scheduling. By applying these concepts, we will explain how we implemented the simple scheduling program and its results for different algorithms, such as first-come-first-served (FCFS), shortest job first (SJF), round-robin (RR), and completely fair safe (CFS). Also, we will show the performance of each algorithm based on the features discussed in a later section. At the end of the paper, we will present the result of the execution of the different scheduling algorithms and compare them.

2 Requirements

Level	Process	Index	Requirement
Basic	Parent	1	Create 10 child processes.
		2	Schedule the child processes according to the Round Robin scheduling policy.
		3	Receive ALARM signal periodically by registering the timer event.
		4	Maintain run-queue and wait-queue.
		5	Accounts the remaining time quantum of all the child process and gives time slice to the child process by sending IPC message through the message queue using system calls.
	Child	1	Workload must consists of infinite loop of dynamic CPU burst and I/O burst.
		2	The values of CPU burst and I/O burst are generated randomly.
		3	When the process receives the time slice from the operating system, it makes the progress.
Optional	Parent	4	While parent process sends the IPC message to the current child process, if the child process decreases the CPU burst value.
		1	If the parent process gets the message from the child process, it checks whether the child process begins I/O burst or not.
		2	If the current process finished the CPU burst, the parent process puts it into the wait queue. If not, the current process is rescheduled again.
	Child	3	For every time tick, parent process decreases the I/O value of the processes in the wait queue.
		1	Child process makes I/O request after CPU burst. Therefore, child process must account the remaining CPU burst.
		2	If the CPU burst reaches to zero, the child sends the IPC message to the parent process with the next I/O burst.

Figure 1 - Requirement Specification

Figure 1 shows the requirements for a simple scheduling program. The implementations for these requirements will be described in detail afterwards.

3 Concepts

3.1 Process




Figure 2 - Layout of the Process in the Memory (Silberschatz et al., 2014)

Modern computers allow multiple programs to be loaded into the memory and executed concurrently. This operation requires firmer control and more compartmentalization of the various programs. To fulfill this requirement, the operating systems need the notion of a process, which is a program in execution. A process is the unit of work in a modern time-sharing system.

The process is an instance of a program in execution. A process is more than the program code, which is known as a text section. It is also consisting of a program counter (PC) and the contents of the registers that present the current state of the process. A process also includes the stack, which contains the temporary data, the data section, which contains global variables, and the heap, which is the memory that is dynamically allocated during the process runtime. The structure of the process is presented in Figure 2.

3.1.2 Process State

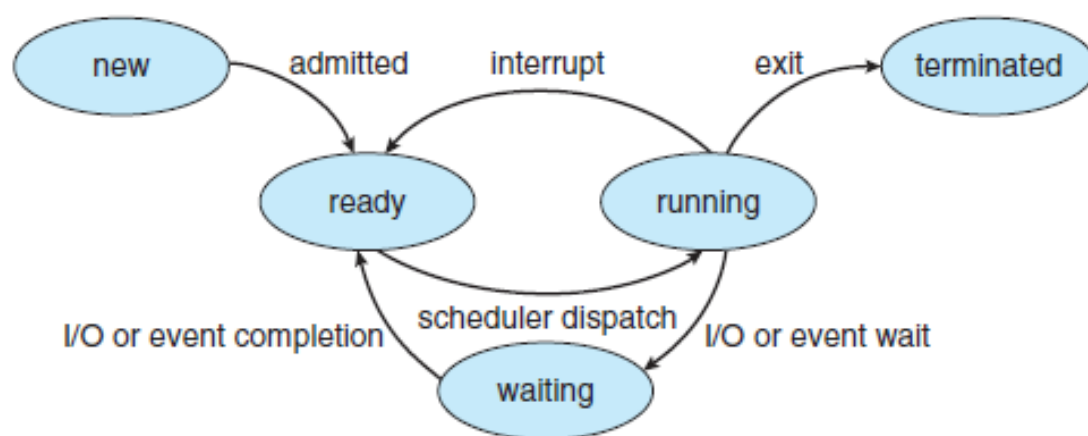


Figure 3 - Diagram of the Process State (Silberschatz et al., 2014)

As a process is in execution, it changes its state. The state of a process can be defined in the part by the current activity of the process. The process may be in one of the following states:

- **New:** The process is being created.
- **Running:** Instructions are being executed.
- **Waiting:** The process is waiting for some event to occur.
- **Ready:** The process is waiting to be assigned to a process.
- **Terminated:** The process has finished the execution.

It is important to realize that only one process can be running on any processor at any instant. However, many processes may be ready and waiting. The state diagram that presents these states is presented in Figure 3.

3.1.3 Process Control Block

process state
process number
program counter
registers
memory limits
list of open files
...

Figure 4 - Process Control Block (PCB) (Silberschatz et al., 2014)

Each process in the operating system is presented in a process control block (PCB), which is also called a task control block. The PCB is presented in Figure 4. It contains much information associated with a specific process, including the following information

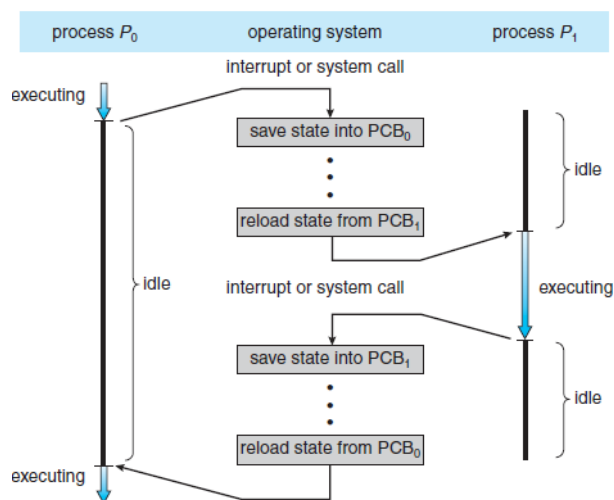


Figure 5 - Diagram of the Context Switching from Process to Process (Silberschatz et al., 2014)

- **Process state:** The state may be new, ready, running, waiting, or terminated.
- **Program counter:** The program counter indicates the address of the next instruction to be executed for this process.
- **CPU registers:** Depending on the computer architecture, the registers vary in number and type. They include accumulators, index registers, stack pointers, and general-purpose registers. Along with the program counter, the following state information must be saved when an interrupt occurs, to allow the process to be continued correctly afterward. The following process is presented in Figure 5.
- **CPU scheduling information:** This information includes a process priority, pointers to scheduling queues, and any other scheduling parameters.
- **Memory management information:** This information may include such items as the value of the base and limit registers and the page table, or the segment tables, depending on the memory system used by the operating system.

- **Accounting information:** This information includes the amount of CPU and real-time used, time limits, account numbers, and job or process numbers.
- **I/O status information:** This information includes the list of I/O devices allocated to the process and a list of open files.

In short, the PCB simply used as the repository for any information that may vary from process to process.

3.2 Process Scheduling

In multiprogramming, it has some process running at all times, to maximize CPU utilization. Also, time-sharing switches the CPU among processes so that the users can interact with each program frequently while it is running. To implement the following operations, the process scheduler selects an available process for program execution on the CPU. In the case of the single-processor system, there will never be more than one running process. If there are more processes, the rest of the processes will have to wait until the CPU is free and can be rescheduled.

3.2.1 Scheduling Queues

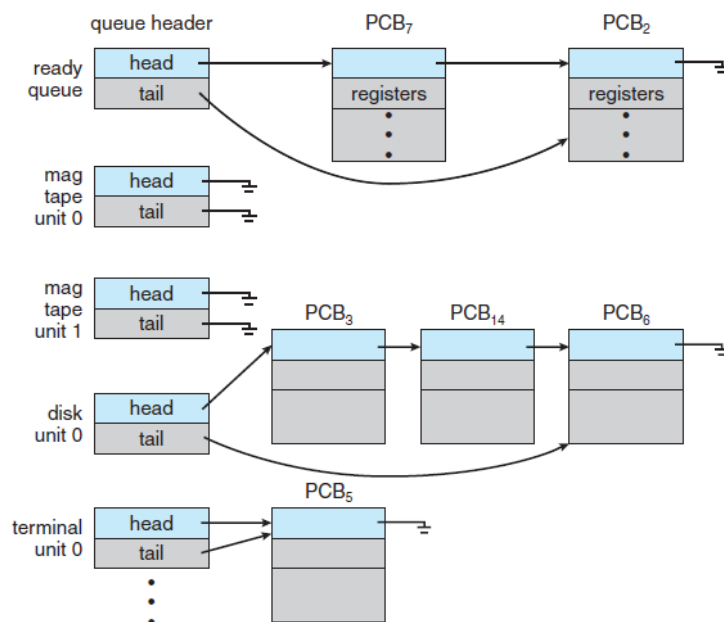


Figure 6 - Multiple Queues in the Operating System (Silberschatz et al., 2014)

As a process enters the system, it is put into the job queue, which consists of all processes in the system. The process that is residing in the main memory and are ready and waiting to execute are kept on a list called the ready queue. The queue is generally stored as a linked list. A ready queue header contains the pointers to the first and final PCBs in the list. Each PCB includes a pointer field that points to the next PCB in the ready queue.

The system also includes the other queues. When a process is allocated the CPU, it executes for a while and eventually quits, is interrupted, or waits for the occurrence of a particular event, such as the completion of an I/O request. Since there are many processes in the system, the disk may be busy with the I/O request of some other process. Therefore, the process may have to wait for the disk. The list of processes waiting for a particular I/O device is called a device queue. The following queue is presented in Figure 6.

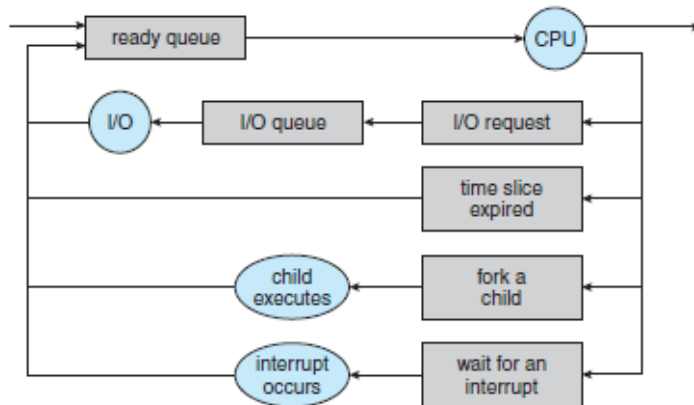


Figure 7 - Queueing-Diagram Representation of Process Scheduling (Silberschatz et al., 2014)

Figure 7 shows the representation of the queuing diagram. Each box in Figure 7 represents a queue. There are two types of queues, which are a ready queue and a set of device queues. The circles represented in Figure 7 are the resources that serve the queues, and the arrow indicates the flow of processes in the system.

A new process is initially put in the ready queue. It waits in the ready queue until it is selected for execution, or dispatched. Once the process is allocated the CPU and is in execution, the following events can occur:

- The process can issue an I/O request and then be placed in an I/O queue.
- The process can create a new child process and wait for the child's termination.
- The process can be removed forcibly from the CPU, as a result of an interrupt, and be put back in the ready queue.

In the first two cases, the process switches from the waiting state to the ready state and is then put back in the ready queue. A process continues this process until it terminates, at the time it is removed from all queues and has its PCB and resources deallocated.

3.2.2 Scheduler

A process migrates among the various scheduling queues throughout its lifetime. The operating system must select, for scheduling purposes, processes from these queues in some method. The selection of the process is carried out by the appropriate scheduler. There are two schedulers, which are long-term scheduler (job scheduler), and the short-term scheduler (CPU scheduler),

The long-term scheduler selects the processes to form the process pool and loads them into the memory for execution. The short-term scheduler selects the process that is ready to execute and allocates the CPU to one of them. The primary distinction between these two schedulers is in the frequency of the execution. The short-term scheduler must select a new process for the CPU frequently. However, the long-term scheduler executes with much less frequency, to create a new process. Therefore, it controls the degree of multiprogramming. If the degree of multiprogramming is stable, the average rate of the process creation must be equal to the average departure rate of the processes leaving the system. Also, the long-term scheduler may need to be invoked only when a process leaves the system. Due to the long interval between the executions, the long-term scheduler can afford to take more time to decide which process should be selected for the execution.

In short, the long-term scheduler must make a careful selection. In most cases, most processes can be described as either I/O bound or CPU bound. An I/O bound process spends more of its time doing I/O operation than it spends doing computation. In contrast, a CPU-bound process generates I/O requests infrequently, using more of its time doing computations. The long-term scheduler must select a good process mix of I/O-bound and CPU-bound processes. If all of the processes are I/O bound, the ready queue will be always empty, and the short-term scheduler will have nothing to do. If all of the processes are CPU bound, the I/O waiting queue will almost always be empty, hardware resources are not used, and again the system will be unbalanced. The system with the best performance will have a combination of CPU-bound and I/O-bound processes.

3.3 Inter-Process Communication (IPC)

Processes executing concurrently in the operating system may be either independent processes or cooperating processes. An independent process cannot affect or be affected by the other processes executing in the system. Any process that does not share data with any other process is independent. However, a process is cooperating if it can affect or be affected by the other processes executing in the system, which means that any process that shared data with other processes is a cooperating process.

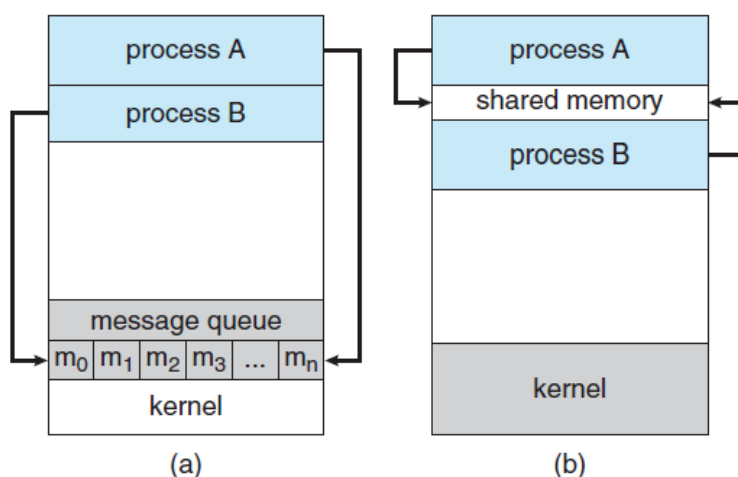


Figure 8 - Communications Models of Shared memory and Message Passing (Silberschatz et al., 2014)

Cooperating processes require an inter-process communication (IPC) operation that will allow them to exchange data and information. There are two fundamental models of inter-process communication, which are shared memory and message passing. In the shared memory model, a region of the memory that is shared by cooperating processes is established. The process can then exchange information by reading and writing data to the shared region. In the message-passing model, communication takes place to be means of messages exchanged between the cooperating processes. The following models are presented in Figure 8.

Both of the models are commonly used in operating systems. Message passing is useful for exchanging smaller amounts of data because no conflicts need to be avoided. Shared memory can be faster than message passing since message-passing systems are typically implemented using system calls and require the more time-consuming task of kernel intervention. In shared memory systems, system calls are required only to establish shared memory regions. Once the shared memory is established, all accesses are treated as routine memory accesses, and no assistance from the kernel is required.

Latest researches indicate that message passing provides better performance than the shared memory on such systems. Shared memory suffers from cache coherency issues, which arise because shared data migrate among several caches. As the number of processing cores on the systems increases, it is possible that message passing is the preferred method for the IPC.

3.3.1 Message Passing




Figure 9 - Message Queue Structure (Silberschatz et al., 2014)

Message passing provides an operation that allows processes to communicate and synchronize their action without sharing the same address space. If processes A and B want to communicate, they must send messages to and receive messages from each other. A communication link must exist between them. This link can be implemented in a variety of ways: direct or indirect communication, synchronous or asynchronous communication, and automatic or explicit buffering.

3.3.2 Message Queue Naming

Processes that want to communicate must have a way to refer to each other. They can use either direct or indirect communication. Under direct communication, each process that wants to communicate must explicitly name the receiver or sender of the communication. This means that both the sender process and receiver process must name the other to communicate.

A variant of this concept is indirect communication. In this case, only the sender names the receiver, while the receiver does not require the name of the sender. With indirect communication, the messages are sent to and received from the mailboxes, also known as ports. A mailbox can be viewed abstractly as an object into which messages can be placed by the processes and from which messages can be removed. Each mailbox has a unique identification. A process can communicate with another process with several different mailboxes, but two processes can communicate only if they have a shared mailbox.

The process that creates a new mailbox becomes the mailbox's owner by default. Initially, the owner is the only process that can receive messages through this mailbox. However, the ownership and receiving privilege may be passed to other processes through appropriate system calls. The following provision can result in multiple receivers for each mailbox.

3.3.3 Message Queue Synchronization

Communication between processes takes place through calls to send and receive primitives. The implementation of these primitives can be differed from blocking to nonblocking, which are also known as synchronous and asynchronous. Each primitive can be represented as below:

- **Blocking send:** The sending process is blocked until the message is received by the receiving process or by the mailbox.
- **Nonblocking send:** The sending process sends the message and resumes the operation.
- **Blocking receive:** The receiver blocks until the message is available.
- **Nonblocking receive:** The receiver retrieves either a valid message or not.

3.3.4 Message Queue Buffering

Generally, the queues can be implemented in three ways:

- **Zero capacity:** The queue has a maximum length of zero, and the link cannot have any messages waiting in it. In the following case, the sender must block the queue until the recipient receives the message.
- **Bounded capacity:** The queue has a finite length n of the queue, and n messages can reside in it. If the queue is not full when a new message is sent, the message is placed in the queue, and the sender can continue execution without waiting. However, the capacity of the link is finite. Therefore, if the link is full, the sender must block it until the space is available in the queue.
- **Unbounded capacity:** The length of the queue is infinite, and any number of messages can wait in the queue. The sender never blocks the queue.

The zero-capacity case is referred to as a message system with no buffering. Other cases are referred to as systems with automatic buffering.

3.4 CPU Scheduling

Whenever the CPU becomes idle, the operating system must select the process in the ready queue to be executed. The selection of the process is carried out by the CPU scheduler. CPU scheduler selects a process from the process in the memory that is ready to execute and allocates the CPU to that process.

We denote that the ready queue is not necessarily a first-in, first-out (FIFO) queue. We can consider the various scheduling algorithms, and the ready queue can be implemented in not only the FIFO queue, but also the priority queue, tree, or simply an unordered linked list. However, all the processes in the ready queue are lined up and waiting for a chance to run on the CPU. The records in the queues are generally the process control block (PCB) of the processes.

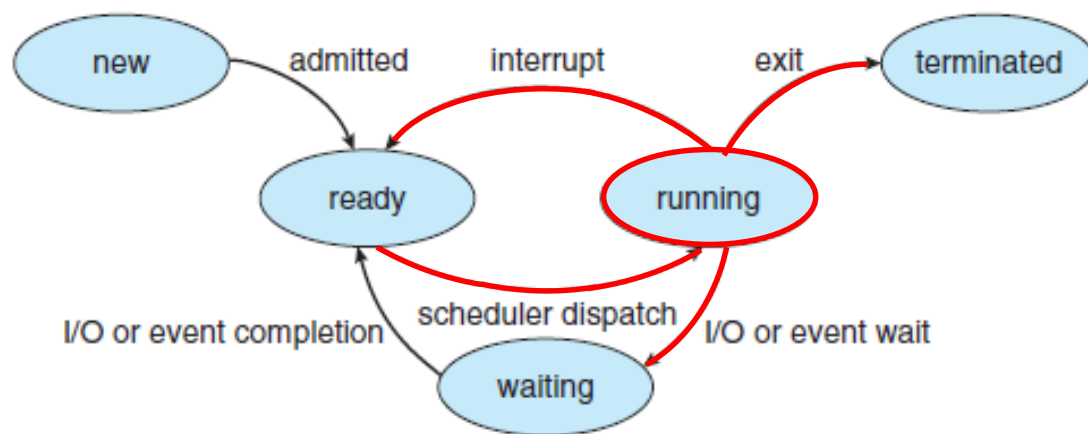


Figure 10 - Diagram of the Process State where Scheduler Affects (Silberschatz et al., 2014)

CPU scheduling decisions can take place under the following circumstances, which are presented in Figure 10:

- When a process switches from a running state to a waiting state.
- When a process switches from a running state to a ready state.
- When a process switches from the waiting state to the ready state.
- When the process terminates.

For situations 1 and 4, there is no choice in terms of scheduling. A new process must be selected for execution. However, there is a choice for situations 2 and 3. When scheduling takes place only under 1 and 4, the scheduling is non-preemptive or cooperative. In contrast, it is preemptive.

Under non-preemptive scheduling, once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by terminating or by switching to the waiting state. However, preemptive scheduling can result in race conditions when the data are shared among several processes.

3.4.1 Scheduling Criteria

Different CPU scheduling algorithms have different properties, and the choice of a particular algorithm may favor one class of processes over another. In choosing which algorithm to use in the situation, we must consider the properties of the various algorithms.

The following criteria are the characteristics of the scheduling algorithms that are used for the comparison to make a substantial difference in which algorithm is judged to be best:

- **CPU utilization:** What we want to do is to keep the CPU as busy as possible. The range that shows how busy the CPU is working is CPU utilization.
- **Throughput:** If the CPU is busy executing three processes, then the work is being done. Throughput presents the number of completed processes per time unit.
- **Turnaround time:** From the point of view of a particular process, the important criterion is how long it takes to execute the process. The interval from the time of the submission of a process to the time of the completion is the turnaround time. This time stands for the sum of periods spent waiting in the ready queue, executing on the CPU, and doing the I/O operations.
- **Waiting time:** The CPU scheduling algorithm does not affect the amount of time during which a process executes or does an I/O operation. It affects only the amount of time that process spends waiting in the ready queue. Waiting time is the sum of the periods spent waiting in the ready queue.
- **Response time:** In an interactive system, turnaround time may not be the best criterion. Often, a process can produce some output early and can continue computing the new results while the previous results are outputting. Also, another measure is the time from the submission of a request until the first response is produced. The following measurement is the response time, which is the time it takes to start responding.

It is desirable to maximize the CPU utilization and throughput and to minimize the turnaround time, waiting time, and response time. In general, we optimize the average measurement. However, under some circumstances, we prefer to optimize the minimum or maximum values rather than the average.

3.4.2 Scheduling Algorithms

CPU scheduling deals with the problem of deciding which of the processes in the ready queue is to be allocated CPU core. There are many different CPU scheduling algorithms. In the later sections, we will describe some of them with the implementation. Although most modern CPU architectures have multiple processing cores, we will describe the scheduling algorithms in the context of a single processing core, which means that the system is only capable of running one process at a time.

4 Simple Scheduling

4.1 Signal and Handler

A program, which is also known as a process, must deal with unexpected or unpredictable events, such as a floating-point error, an alarm clock rings, a termination request from a user, and so on. In the following situations, the signal is used. Signals are the software interrupts and provide a way to handle asynchronous events. Once a signal is generated by some event, it is delivered to a process, which then takes some action in response to the signal. Upon delivery of a signal, a process can operate in the following ways:

- **Ignore it:** literally, ignore the signal.
- **Catch it:** A programmer can arrange for a particular signal to be processed by a special piece of code called a signal handler.
- **Accept the default:** Have the kernel do whatever is defined as the default action for the signals, such as the signal is ignored, the process is terminated, the process is stopped, and so on.

SIGNAL	NBR	DESCRIPTION	DEFAULT ACTION
SIGABRT	6	Process abort (sent by assert)	Implementation dependent
SIGALRM	14	alarm clock	Abnormal termination
SIGCONT	25	Execution continued if stopped	Continue
SIGFPE	8	Arithmetic error (i.e., divide by zero)	Implementation dependent
SIGINT	2	Interactive attention signal (Ctrl-C)	Abnormal termination
SIGKILL	9	Termination (cannot be caught or ignored)	Abnormal termination
SIGQUIT	3	Interactive termination; core dump	Implementation dependent
SIGSEGV	11	Segmentation fault	Abnormal termination
SIGSTOP	23	Execution stop (cannot be caught or ignored)	Stop
SIGTERM	15	Termination	Abnormal termination
SIGUSR1	16	User-defined signal 1	Abnormal termination
SIGUSR2	17	User-defined signal 2	Abnormal termination ⁷

Figure 11 - POSIX Signals

In the signal header file, which is offered by the POSIX system, the signals presented in Figure 11 are defined as the integer value.

Signals can occur during the process execution, and we can also generate the signal explicitly to another process. Generating the signal explicitly is called signal dispatching.

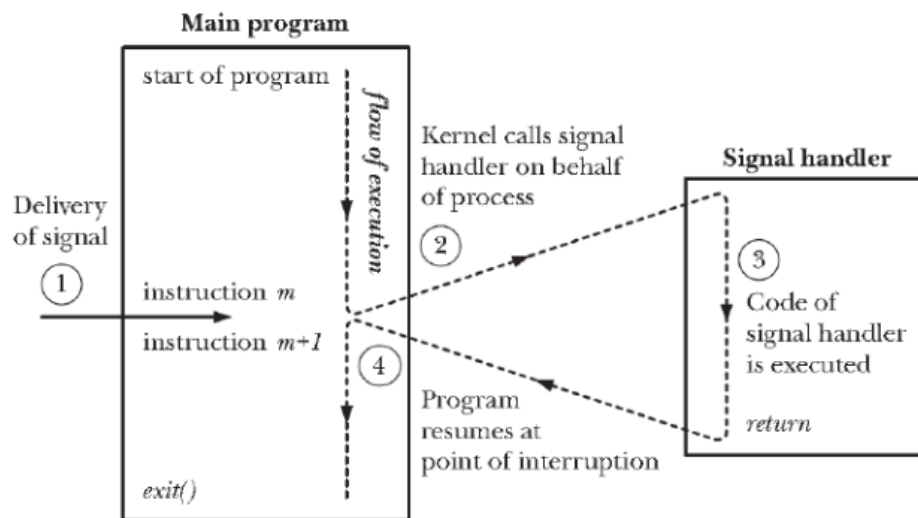


Figure 12 - Signal Delivery and Handler Execution

As the signals are generated asynchronously, we cannot predict the period of the signal generation, which means that the signal can be generated whenever the program is in execution. Figure 12 presents the operation of the signal delivery and the handler execution. According to Figure 12, first, if the signal occurs while the program is executing, it temporarily stops the execution. Next, the signal handler is executed to proceed with the signal. Then, right after the signal proceeds, the next instruction, right after the signal is executed.

```
struct sigaction {
    void (*sa_handler)(int);
    void (*sa_sigaction)(int, siginfo_t *, void *);
    sigset_t sa_mask;
    int sa_flags;
    void (*sa_restorer)(void);
};
```

Figure 13 – Structure of the SIGACTION System Call

The POSIX system provides the sigaction system call to register the signal handler. The following system call is used to change the action taken by a process on receipt of a specific signal. Figure 13 presents the structure of the sigaction system call.

In this project, we set three main signals, which are the I/O signal, CPU signal, and time clock (tick) signal. In the I/O signal, the program ferrets out each PCB in the waiting queue and decreases the I/O burst time. If the I/O burst time is zero, it pops the PCB from the waiting queue and appends it into the ready queue. In the CPU signal, it selects the next executing process from the ready queue and decreases the CPU burst time. In the time clock (tick) signal, it increases the time tick of the CPU, and prints out the result after executing the I/O and CPU signal.

4.2 Message Passing in System V Messages

System V messages allow the processes to exchange data in the form of messages. Messages can be assigned to a specific type. The server process can direct message traffic between clients to its queue by using the client process PID as the message type. For single message transactions, multiple server processes can work in parallel with transactions that are sent to a shared message queue.

Before a process can send or receive a message, it must initialize the queue through the `msgget` system call. The owner or creator of the queue can change their ownership or permissions using the `msgctl` system call. Any process with permission can use the `msgctl` system call for control operation. Operations to send and receive messages are performed by `msgsnd` and `msgrcv` system calls. When a message is sent, its text is copied to the message queue. `msgsnd` and `msgrcv` system calls can be performed as either blocking or non-blocking operations. The blocked message operation remains suspended until one of the following three conditions occurs: the call succeeds, the process receives a signal, and the queue is removed.

```
#include <sys/ipc.h>
#include <sys/msg.h>

...
key_t   key;           /* key to be passed to msgget() */
int     msgflg, /* msgflg to be passed to msgget() */
        msqid; /* return value from msgget() */

...
key = ...
msgflg = ...
if ((msqid = msgget(key, msgflg)) == -1)
{
    perror("msgget: msgget failed");
    exit(1);
} else
    (void) fprintf(stderr, "msgget succeeded");
...
```

Figure 14 - msgget System Call Code

`msgget` system call initializes a new message queue. It can also return the message queue ID to the queue corresponding to the key argument. The values passed as the message flag argument must be an octal integer with a setting for the queue's permissions and control flags. The following system call is presented in Figure 14.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

...
if (msgctl(msqid, IPC_STAT, &buf) == -1) {
    perror("msgctl: msgctl failed");
    exit(1);
}

...
if (msgctl(msqid, IPC_SET, &buf) == -1) {
    perror("msgctl: msgctl failed");
    exit(1);
}

...
```

Figure 15 - msgctl System Call Code

msgctl system call alters the permissions and other characteristics of the message queue. The message queue ID argument must be the ID of an existing message queue. The commands of the msgctl system call are the IPC_STAT (place information about the status of the queue in the data structure pointed to buffer), the IPC_SET (Set the owner's user and group ID, the permissions, and the size of the message queue), and the IPC_RMID (remove the message queue specified by the message queue ID argument). The following system call is presented in Figure 15.

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

...
int          msgflg; /* message flags for the operation */
struct msgbuf *msgp; /* pointer to the message buffer */
size_t       msgsz;  /* message size */
size_t       maxmsgsize;
long         msgtyp; /* desired message type */
int          msqid    /* message queue ID to be used */
...
msgp = malloc(sizeof(struct msgbuf) - sizeof (msgp->mtext)
              + maxmsgsz);
if (msgp == NULL) {
    (void) fprintf(stderr, "msgop: %s %ld byte messages.\n",
                  "could not allocate message buffer for", maxmsgsz);
    exit(1);
    ...
    msgsz = ...
    msgflg = ...
    if (msgsnd(msqid, msgp, msgsz, msgflg) == -1)
        perror("msgop: msgsnd failed");
    ...
    msgsz = ...
    msgtyp = first_on_queue;
    msgflg = ...
    if (rtrn = msgrcv(msqid, msgp, msgsz, msgtyp, msgflg) == -1)
        perror("msgop: msgrcv failed");
    ...
}
```

Figure 16 - msgsnd and msgrcv System Calls Code

msgsnd and msgrcv system calls send and receive messages. The message queue ID argument must be the ID of the existing message queue, The message pointer argument is a pointer to a structure that contains the type of the message and its text. The message size argument specifies the length of the message in bytes. The message flag argument passes the various control flags. The following system call is presented in Figure 16.

4.3 First-Come First-Served Algorithm (FCFS)

The simplest CPU scheduling is the first-come, first-serve (FCFS) scheduling algorithm. In this algorithm, the process that requests the CPU resources first is allocated to the CPU first. The implementation of the FCFS policy is simply managed by a FIFO queue. When a process enters the ready queue, its PCB is linked to the tail of the queue. The running process is then removed from the queue. The code for FCFS scheduling is simple to write and understand.

FCFS (Example)

Process	Duration	Oder	Arrival Time
P1	24	1	0
P2	3	2	0
P3	4	3	0

Gantt Chart :



P1 waiting time : 0

P2 waiting time : 24

P3 waiting time : 27

The Average waiting time :

$$(0+24+27)/3 = 17$$

Figure 17 - Example of FCFS Scheduling

Figure 17 shows an example of the FCFS scheduling case. As we can see from Figure 17, even though the FCFS is easy to implement, it has the disadvantage that its average waiting time is often quite long. Also, consider the performance of FCFS scheduling in a dynamic situation. If we have one CPU-bound process and multiple I/O-bound processes. As the process flows around the system, the CPU-bound process will get and hold the CPU. During this time, all other processes will finish their I/O and will move into the ready queue, waiting for the CPU. While the processes wait in the ready queue, the I/O devices are idle. After the CPU-bound process finishes its CPU bursts and moves to an I/O device, all the I/O-bound processes, which have short CPU bursts, execute quickly and move back to the I/O queues. At this point, the CPU sits idle. The following situation is called the convoy effect, which presents the case of all the other processes waiting for the one big process to get off the CPU. This effect results in lower CPU and device utilization than might be possible if the shorter processes were allowed to go first.

4.4 Shortest Job First Algorithm (SJF)

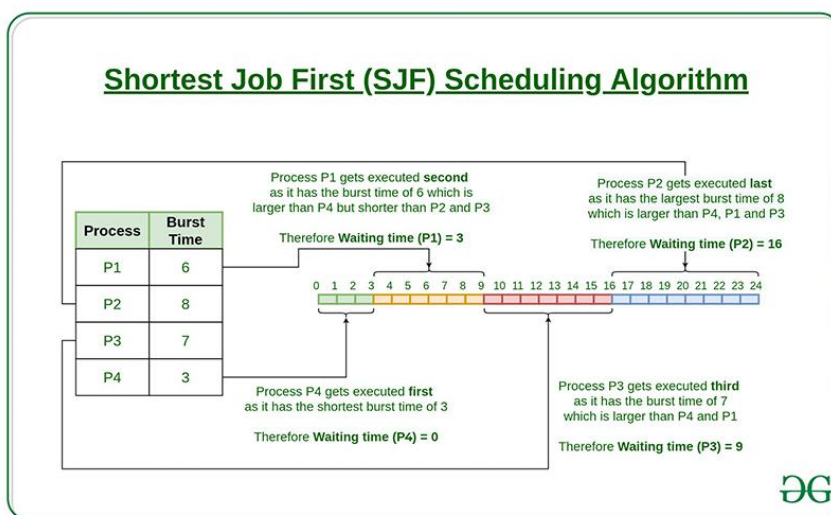


Figure 18 – Example of SJF Scheduling

A different approach to CPU scheduling is the shortest job first scheduling algorithm. This algorithm is associated with each process, the length of the process's next CPU burst time. When CPU resources are available, it is assigned to the process that has the smallest next CPU burst time. If the next CPU bursts of the two processes are the same, FCFS scheduling is used. Figure 18 presents the execution of the SJF scheduling.

The SJF scheduling algorithm is probably optimal, due to that gives the minimum average waiting time for a given set of processes. Moving a short process before a long one decreases the waiting time of the short process more than it increases the waiting time of the long process. As a result, the average waiting time decreases.

However, the SJF scheduling algorithm cannot be implemented at the level of CPU scheduling, as there is no way to know the length of the next CPU burst. One approach to this problem is to try to approximate SJF scheduling. We may not know the length of the next CPU burst, but we may be able to predict its value. By computing an approximation of the length of the next CPU burst, we can pick the process with the shortest predicted CPU burst.

4.5 Round Robin Algorithm (RR)

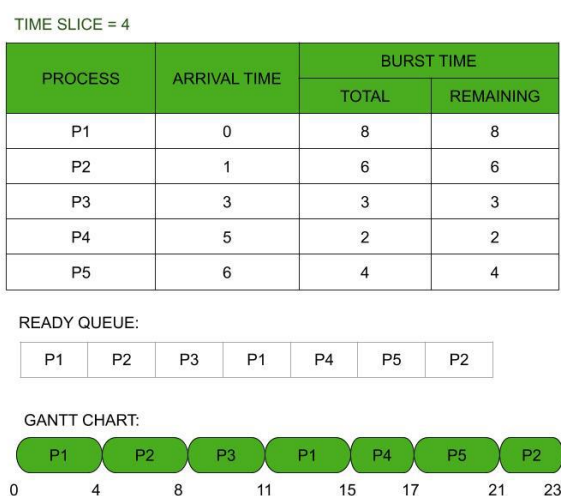


Figure 19 - Example of RR Scheduling

The round-robin scheduling algorithm is like FCFS scheduling, but preemption is added to enable the system to switch between processes. A small unit of time, which is time quantum or time slice, is defined. Generally, time quantum has a range of 10 to 100 milliseconds in length. The ready queue is treated as a circular queue. The CPU scheduler goes around the ready queue, allocating the CPU to each process for a time interval of up to 1-time quantum. To implement RR scheduling, we need to treat the ready queue as a FIFO queue of processes. New processes have been added to the tail of the ready queue. The CPU scheduler picks the first process from the ready queue, sets a timer to interrupt after 1-time quantum, and dispatches the process. Figure 19 shows the execution of the RR scheduling algorithm.

One of two things will happen when we apply the RR scheduling algorithm. The process may have a CPU burst of less than 1-time quantum. In this case, the process itself will release the CPU voluntarily. However, if the CPU burst of the currently running process is longer than the 1-time quantum, the timer will go off and will cause an interruption to the operating system. Then a context switch will be executed, and the process will be put at the tail of the ready queue.

Therefore, we can say that the performance of the RR scheduling algorithm depends heavily on the size of the time quantum. In one extreme case, if the time quantum is extremely large, the RR policy is the same as the FCFS policy. In contrast, if the time quantum is extremely small, the RR policy can result in many context switches.

Also, turnaround time depends on the size of the time quantum. The average turnaround time of a set of processes does not necessarily improve as the time quantum increases. Generally, the average turnaround time can be improved if most processes finish their next CPU burst in a single time quantum. If context switch time is added in, the average turnaround time increases even more for a smaller time quantum, as more context switches are required.

4.6 Program Definition

Before implementing the multi-threaded word count program, we will state the additional program definition that will be used in the real implementation.

- Global Variables

Variables	Data Type	Definition
RUN_TIME	10000	Limitation of the program run time.
MAX_PROCESS	10	Number of the child processes.
TIME_QUANTUM	5	Time quantum for the round robin (RR) algorithm.
TIME_TICK	1000	SIGALRM interval time, which is 0.001 second (10 millisecond).
Counter	Int	Counts the time tick of CPU.
Run_time	Int	Substitution of the RUN_TIME and the Counter.
Readyq	Void*	Pointer that points at the ready queue.
Waitq	Queue*	Pointer that points at the waiting queue.
Max_limit	Int	Limitation of the maximum CPU and I/O burst time.
Proc_count	Int	Counter that counts the completed process.
Cur_wait	Node*	Pointer that points at the current node in the waiting queue.
Cur_ready	Node*	Pointer that points at the current node in the ready queue.
Wait_time	Double	Variable used to calculate the total waiting time among the processes.
Set_scheduler	Int	Index that is used to choose the scheduling algorithm.
Turnaround_time	Double	Variable used to calculate the total turnaround time among the processes.
Key	Int[]	Keys that are used to access the message queue between processes.
Cpid	Pid_t[]	Array for storing process ID of the child processes.
Io_burst	Int[]	Array for storing I/O burst time of the child processes.
Cpu_burst	Int[]	Array for storing CPU burst time of the child processes.
End_time	Int[]	Array for storing the end time of the child processes.
Service_time	Int[]	Array for storing the service time of the child processes.
Arrival_time	Int[]	Array for storing the arrival time of the child processes.

- Modules and Functions

Modules	Functions	Definition
Queue	CreateNode	Creates the new node and returns it.
	IsEmpty	Returns whether the following queue is empty or not.
	CreateQueue	Creates the new queue and returns it.
	Dequeue	Dequeues the node from the front of the following queue.
	RemoveQueue	Removes all the nodes in the queue. After removing them, it frees the queue.
	PrintQueue	Prints all the data for each node in the following queue.
	FprintQueue	Prints all the data for each node in the following queue to the text file.
	Enqueue	Enqueues the new node into the rear of the following queue.
Heap	CreateHeap	Creates the new heap and returns it.
	RemoveHeap	Removes all the nodes in the heap. After removing them, it frees the heap.
	DeleteHeap	Deletes the node from the front of the following heap.
	PrintHeap	Prints all the data for each node in the following heap.
	Heapify	Readjust the following array into the heap.
	FprintHeap	Prints all the data for each node in the following heap to the text file.
	InsertHeap	Insert the new node into the following heap.
Msg	Cmsgsnd	Message send function of the child process. Child process sends the message structure to the parent process through the IPC communication.
	Pmsgrcv	Message receive function of the parent process. Parent process receives the message structure from the child process through the IPC communication.
Main	Dump_data	Writes the data of ready queue dump and wait queue dump on the text file.
	Signal_io	Signal handler that checks whether the child process has the IO burst or not.
	Signal_rr	Signal handler that enqueues the current CPU process into the end of the ready queue and dequeues the next process from the ready queue which is to be executed. The applied scheduling algorithm is round robin policy.
	Signal_sjf	Signal handler that enqueues the current CPU process into the end of the ready queue and dequeues the next process from the ready queue which is to be executed. The applied scheduling algorithm is shortest job first policy.
	Signal_fcfs	Signal handler that enqueues the current CPU process into the end of the ready queue and dequeues the next process from the ready queue which is to be executed. The applied scheduling algorithm is first-come, first-served policy.
	Signal_count	Signal handler of SIGALRM which is called for every count.

5 Implementation

5.1 Queue Header

```

typedef struct PCB {
    int idx;
    pid_t pid;
    int io_burst;
    int cpu_burst;
} PCB;

typedef struct Node {
    PCB pcb;
    struct Node* next;
} Node;

typedef struct Queue {
    int count;
    Node* tail;
    Node* head;
} Queue;

```

Figure 20 - Structures used in the Queue

Figure 20 shows the process control block (PCB), node, and queue structures that are used in the implementation of the queue. As the PCB stores the information of the process, we set the PCB to contain the index, process ID (PID), I/O burst time, and CPU burst time. Then, in the node, which is used in the queue, we set the node to contain the PCB with the node pointer that points to the next node in the queue. Finally, in the queue structure, we set it to contain the count, which represents the number of the node in the queue, tail, and head, which are pointing to the last node and the first node of the queue.

```
Node* createNode();
int isEmpty(Queue* q);
Queue* createQueue();
Node* dequeue(Queue* q);
void removeQueue(Queue* q);
void printQueue(Queue* q, char c);
void fprintfQueue(Queue* q, char c, FILE* fp);
void enqueue(Queue* q, int idx, int cpu_burst, int io_burst);
```

Figure 21 - Functions used in the Queue

Figure 21 shows the functions of the queue operations. We can create the node and the queue by using the createNode and createQueue functions. By using dequeue and enqueue functions, we can pop out the first node in the queue and insert the node into the end of the queue structure. Each printQueue and fprintfQueue function are used to output the progress and the results of the program execution on the screen and the dump file. After all the operations, we can remove the queue structure and the remaining nodes by using the removeQueue function.

The following queue header is used in the implementations of first-come, first-served (FCFS) and round-robin (RR) scheduling algorithms.

5.2 Heap Header

```
typedef struct Heap {
    int count;
    Node** heap;
} Heap;
```

Figure 22 - Structure used in the Heap

Figure 22 shows the heap structures that are used in the implementation of the heap, which is also called a priority queue. In the case of the process control block (PCB) and node structure, we used the same ones that we implemented in the queue header. In the heap structure, we set it to contain the count, which represents the number of nodes in the heap, and the heap pointer, which is pointing the root node of the heap.

```

Heap* createHeap();
void removeHeap(Heap* h);
Node* deleteHeap(Heap* h);
void printHeap(Heap* h, char c);
void heapify(Heap* h, int index);
void fprintfHeap(Heap* h, char c, FILE* fp);
void insertHeap(Heap* h, int idx, int cpu_burst, int io_burst);

```

Figure 23 - Functions used in the Heap

Figure 23 shows the functions of the heap operations. We can create the heap by using the createHeap function. By using deleteHeap and insertHeap functions, we can delete and get the root node in the heap and insert the node into the end of the heap structure and perform the heapify operation. Each printHeap and fprintfHeap function are used to output the progress and the results of the program execution on the screen and the dump file. After all the operations, we can remove the heap structure and the remaining nodes by using the removeHeap function.

The following queue header is used in the implementation of shortest job first (SJF) scheduling algorithms.

5.3 Inter-Process Control (IPC) Message Passing Header

```

typedef struct msgbuf {
    long mtype;

    // pid will sleep for io_time
    int pid;
    int io_time;
    int cpu_time;
} msgbuf;

```

Figure 24 - Structure of the Message Buffer

According to the previous section, we can see that inter-process control (IPC) message passing is one of the operations that can control from one process to another process by sending and receiving the messages. To implement the following operations, we implemented the structure of the message that is used in the message-passing communication between the parent process and child processes. First, we defined the type of the message, to specify the role of the message. Then, we set the message structure to contain the information of process id (PID), I/O burst time, and CPU burst time, to notice that the child process has processed how much time tick for its CPU or I/O burst operation. After getting this message from the child processes, the parent process can identify what has happened to the child processes during the time tick. The implemented code of the message structure is presented in Figure 24.

```

void cmsgsnd(int k, int c, int io) {
    int qid = msgget(k, IPC_CREAT | 0666);

    msgbuf msg;
    memset(&msg, 0, sizeof(msgbuf));

    msg.mtype = 1;
    msg.io_time = io;
    msg.cpu_time = c;
    msg.pid = getpid();

    if (msgsnd(qid, &msg, sizeof(msgbuf) - sizeof(long), 0) == -1) {
        perror("msgsnd");
        exit(EXIT_FAILURE);
    }
}

```

Figure 25 - Message Send Function of the Child Process

Figure 25 shows the message send function that is used in the child process. First, the child process gets the message queue ID (QID) from the message queue. Then by setting the parameters in the message structure, it sends the message to the parent process through the message queue that has the ID of QID. If the message is sent successfully, it terminates the operation. If not, it will print out the error statement.

```

void pmsgrcv(int idx, Node* node) {
    int k = 0x12345 * (idx + 1);
    int qid = msgget(k, IPC_CREAT | 0666);

    msgbuf msg;
    memset(&msg, 0, sizeof(msgbuf));

    if (msgrcv(qid, &msg, sizeof(msgbuf) - sizeof(long), 0, 0) == -1) {
        perror("msgrcv");
        exit(EXIT_FAILURE);
    }

    node->pcb.pid = msg.pid;
    node->pcb.io_burst = msg.io_time;
    node->pcb.cpu_burst = msg.cpu_time;
}

```

Figure 26 - Message Receive Function of the Parent Process

Figure 26 shows the message-receiving function that is used in the parent process. First, the parent process generates the key that is used to create the message queue and gets the message queue ID (QID) by using the following key. After that, the parent process checks whether the message has been received in the message queue. If there is a received message in the queue, it reads it. If not, wait for the message to be filled out, and if the error occurs, it will print out the error statement.

5.4 Main

```
void dump_data(FILE* fp) {
    switch (set_scheduler) {
        case 1: fp = fopen("fcfs_dump.txt", "a+"); break;
        case 2: fp = fopen("sjf_dump.txt", "a+"); break;
        case 3: fp = fopen("rr_dump.txt", "a+"); break;
        default:
            perror("scheduler");
            exit(EXIT_FAILURE);
    }

    fprintf(fp, "-----\n");
    fprintf(fp, "Time: %d\n", RUN_TIME - run_time);
    fprintf(fp, "CPU Process: %d -> %d\n", cur_ready->pcb.idx, cur_ready->pcb.cpu_burst);

    for (int i = 0; i < waitq->count; i++) {
        cur_wait = dequeue(waitq);
        fprintf(fp, "I/O Process: %d -> %d\n", cur_wait->pcb.idx, cur_wait->pcb.io_burst);
        enqueue(waitq, cur_wait->pcb.idx, cur_wait->pcb.cpu_burst, cur_wait->pcb.io_burst);
    }

    if (set_scheduler == 2) fprintfHeap(readyq, 'r', fp);
    else fprintfQueue(readyq, 'r', fp);
    fprintfQueue(waitq, 'w', fp);
    fprintf(fp, "-----\n");
    fclose(fp);
}
```

Figure 27 - Dump Data Function

Figure 27 shows the code of the dump data function that dumps the progress of the program execution and the result into the dump file. In this function, we output the status of the ready and waiting queue to the screen and the dump file.

```
/*
 * void signal_io(int signo)
 *
 * Signal handler that checks whether the child process has the IO burst or not.
 * If it has the remained IO burst, it enqueues the process into the wait queue.
 * If not, it enqueues the process into the ready queue.
 */
void signal_io(int signo) {
    pmsgrcv(cur_ready->pcb.idx, cur_ready);
    if (cur_ready->pcb.io_burst == 0) {
        if (set_scheduler == 2) insertHeap(readyq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);
        else enqueue(readyq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);
    }

    else enqueue(waitq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);

    if (set_scheduler == 2) cur_ready = deleteHeap(readyq);
    else cur_ready = dequeue(readyq);
    counter = 0;
}
```

Figure 28 - I/O Operation Signal

Figure 28 presents the implemented code of the I/O operation signal. Ever since the I/O operation signal is generated, it checks the message queue whether there is a message or not. If the received message contains the information of I/O burst time with zero, then it enqueues or inserts it into the ready queue. If not, it enqueues the node into the waiting queue again. After the following progress, it pops out the node in front of the ready queue.

```

/*
 * void signal_fcfs(int signo)
 *
 * Signal handler that enqueues the current cpu process into the end of the ready queue
 * and dequeues the next process from the ready queue which is to be executed.
 */
void signal_fcfs(int signo) {
    cur_ready->pcb.cpu_burst--;
    if (cur_ready->pcb.cpu_burst == 0) {
        enqueue(readyq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);
        cur_ready = dequeue(readyq);
    }
}

/*
 * void signal_sjf(int signo)
 *
 * Signal handler that insert the current cpu process into the end of the ready queue
 * and deletes the next process from the ready queue which is to be executed.
 */
void signal_sjf(int signo) {
    counter++;
    cur_ready->pcb.cpu_burst--;
    if (cur_ready->pcb.cpu_burst == 0) {
        insertHeap(readyq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);
        cur_ready = deleteHeap(readyq);
    }
}

/*
 * void signal_rr(int signo)
 *
 * Signal handler that enqueues the current cpu process into the end of the ready queue
 * and dequeues the next process from the ready queue which is to be executed.
 */
void signal_rr(int signo) {
    counter++;
    cur_ready->pcb.cpu_burst--;
    if (counter >= TIME_QUANTUM) {
        enqueue(readyq, cur_ready->pcb.idx, cur_ready->pcb.cpu_burst, cur_ready->pcb.io_burst);
        cur_ready = dequeue(readyq);
        counter = 0;
    }
}

```

Figure 29 - Signals for CPU Scheduling Algorithms

Figure 29 shows the implemented code of the signals that are used to schedule the CPU with different scheduling policies. First, in a signal of FCFS policy decrease the CPU burst time of the current node, which is the front node of the ready queue, and if the CPU burst time converges to zero, it enqueues the node into the waiting queue and pops out the node from the ready queue. Next, in a signal of SJF policy decreases the CPU burst time of the current node, which is the root node of the ready heap, and if the CPU burst time converges to zero, it enqueues the node into the waiting queue and deletes the node from the ready heap. Third, a signal of RR policy decreases the CPU burst time of the current node, which is the front node of the ready queue, and if the CPU burst time converges to zero, it enqueues the node into the waiting queue. If the CPU burst time does not converge to zero and it spent all its allocated time quantum, the node is enqueued in the ready queue.

```

if (cur_ready->pcb.cpu_burst == 1) {
    proc_count++;
    end_time[cur_ready->pcb.idx] = RUN_TIME - run_time;
    wait_time += end_time[cur_ready->pcb.idx] - arrival_time[cur_ready->pcb.idx];
    turnaround_time += end_time[cur_ready->pcb.idx] - service_time[cur_ready->pcb.idx] - arrival_time[cur_ready->pcb.idx];
}

int length = waitq->count;
for (int i = 0; i < length; i++) {
    cur_wait = dequeue(waitq);
    cur_wait->pcb.io_burst--;
    printf("I/O Process: %d -> %d\n", cur_wait->pcb.idx, cur_wait->pcb.io_burst);
    if (cur_wait->pcb.io_burst == 0) {
        if (set_scheduler == 2) insertHeap(readyq, cur_wait->pcb.idx, cur_wait->pcb.cpu_burst, cur_wait->pcb.io_burst);
        else enqueue(readyq, cur_wait->pcb.idx, cur_wait->pcb.cpu_burst, cur_wait->pcb.io_burst);
        arrival_time[cur_wait->pcb.idx] = RUN_TIME - run_time;
        service_time[cur_wait->pcb.idx] = 0;
        end_time[cur_wait->pcb.idx] = 0;
    }
    else enqueue(waitq, cur_wait->pcb.idx, cur_wait->pcb.cpu_burst, cur_wait->pcb.io_burst);
}
if (set_scheduler == 2) printHeap(readyq, 'r');
else printQueue(readyq, 'r');
printQueue(waitq, 'w');
printf("-----\n");

dump_data(fp);

if (cur_ready->pcb.idx != -1) kill(cpid[cur_ready->pcb.idx], SIGCONT);

if (run_time != 0) {
    run_time--;
}

```

Figure 30 - Signal for Count which is Executed for Every Time Tick

Figure 30 presents the implemented code for the count signal, which is called at every time tick. For every call of the count signal, it performs the following sequence of operations:

1. Check if the current node, which is the front node of the ready queue, is done processing its CPU burst time. If so, the function will record the scheduling criteria, which are the number of completed processes, end time, wait time, and turnaround time.

2. Then, the function will decrease the I/O burst time of every node in the waiting queue. If there is a node with an I/O burst time of zero, the following node is enqueued in the ready queue.
3. After the following progress, the function will dump all the logs and the results on the screen and in the dump file.

By processing the following sequence of operations by calling the count signal, we can simulate the CPU scheduling with the proper scheduling policies.

```
// set the timer for SIGALRM
struct itimerval new_itimer;
struct itimerval old_itimer;

new_itimer.it_interval.tv_sec = 0;
new_itimer.it_interval.tv_usec = TIME_TICK;
new_itimer.it_value.tv_sec = 1;
new_itimer.it_value.tv_usec = 0;

struct sigaction io;
struct sigaction count;
struct sigaction scheduler;

memset(&io, 0, sizeof(io));
memset(&count, 0, sizeof(count));
memset(&scheduler, 0, sizeof(scheduler));

io.sa_handler = &signal_io;
count.sa_handler = &signal_count;

switch (set_scheduler) {
case 1: scheduler.sa_handler = &signal_fcfs; break;
case 2: scheduler.sa_handler = &signal_sjf; break;
case 3: scheduler.sa_handler = &signal_rr; break;
default:
    perror("scheduler");
    exit(EXIT_FAILURE);
}

sigaction(SIGUSR2, &io, NULL);
sigaction(SIGALRM, &count, NULL);
sigaction(SIGUSR1, &scheduler, NULL);
```

Figure 31 - Main Function Part 1

Figure 31 shows the first substantial implemented code of the main function. It presents the initialization of the signals and the timer set of the SIGALRM signal. First, we set the timer for the SIGALRM to 0.001 seconds (10 milliseconds) of timer interval and set the timer to start after the 1 second of the program execution. Next, we initialized the signals for I/O, count, and CPU scheduler, which are presented previously. Then, by using the sigaction function, we generated all three signals, which are called for every timer interval.

```
// create the nodes and the queues for the scheduling operation
waitq = createQueue();
cur_wait = createNode();
cur_ready = createNode();
if (set_scheduler == 2) readyq = createHeap(MAX_PROCESS); else readyq = createQueue();
if (waitq == NULL || readyq == NULL || cur_wait == NULL || cur_ready == NULL) {
    perror("malloc");
    exit(EXIT_FAILURE);
}

// create the message queues for each child process.
for (int i = 0; i < MAX_PROCESS; i++) {
    key[i] = 0x12345 * (i + 1);
    msgctl(msgget(key[i], IPC_CREAT | 0666), IPC_RMID, NULL);
}
```

Figure 32 - Main Function Part 2

Figure 32 shows the second substantial implemented code of the main function. It presents the initialization of the ready queue, waiting queue, current wait node, current ready node, and the messages queues that are used between parent process and child processes.

```

// creates the child process and operates the processes.
for (int i = 0; i < MAX_PROCESS; i++) {
    pid_t pid;

    // error on the fork operation.
    if ((pid = fork()) < 0) {
        perror("fork");
        exit(EXIT_FAILURE);
    }

    // child process
    else if (pid == 0) {
        int idx = i;
        int c_io = io_burst[i];
        int c_cpu = cpu_burst[i];

        // child process waits until the interrupt occurs.
        kill(getpid(), SIGSTOP);

        // child operates the cpu burst.
        while (1) {
            c_cpu--; // decreases the cpu burst by 1.i
            if (c_cpu == 0) {
                c_cpu = cpu_burst[i];

                // send the PCB data of the child process to the parent process.
                cmsgsnd(key[idx], c_cpu, c_io);
                c_io = io_burst[i];
                kill(ppid, SIGUSR2);
            }
            else { // cpu burst is over.
                kill(ppid, SIGUSR1);
            }

            // child proces waits for the next interrupt.
            kill(getpid(), SIGSTOP);
        }
    }

    else {
        cpid[i] = pid;
        arrival_time[i] = RUN_TIME - run_time;
        if (set_scheduler == 2) insertHeap(readyq, i, cpu_burst[i], io_burst[i]);
        else enqueue(readyq, i, cpu_burst[i], io_burst[i]);
    }
}
}

```

Figure 33 - Main Function Part 3

Figure 4 shows the solution code that uses mutex lock to resolve the critical section problem. The acquire function acquires the lock, and the release function releases the lock.

6 Build Environment

- Build Environment:
 1. Linux Environment -> Vi editor, GCC Compiler
 2. Program is built by using the Makefile.
- Build Command:
 1. `$make main` -> build the execution program for simple scheduler.
 2. `$make clean` -> clean all the object files that consists of the simple scheduler programs.
- Execution Command:
 1. `./main 1 {$max_limit}`
-> Execute the simple scheduler program with the scheduler of FCFS policy with the maximum CPU and I/O burst time of {\$max_limit} of the child processes.
 2. `./main 2 {$max_limit}`
-> Execute the simple scheduler program with the scheduler of SJF policy with the maximum CPU and I/O burst time of {\$max_limit} of the child processes.
 3. `./main 3 {$max_limit}`
-> Execute the simple scheduler program with the scheduler of RR policy with the maximum CPU and I/O burst time of {\$max_limit} of the child processes.

7 Results

- Result of the simple scheduler with FCFS policy and 10 maximum burst time:

```
-----
Process 0    -> CPU: 3,    10:    6    Time: 9996
Process 1    -> CPU: 3,    10:    1    CPU Process: 4 -> 2
Process 2    -> CPU: 10,   10:    3    Ready: 5 6 7 8 1 9 0 2 3
Process 3    -> CPU: 4,    10:    7    Wait:
Process 4    -> CPU: 10,   10:    9
Process 5    -> CPU: 7,    10:    3
Process 6    -> CPU: 3,    10:    3    Time: 9997
Process 7    -> CPU: 2,    10:    5    CPU Process: 4 -> 1
Process 8    -> CPU: 10,   10:    2    Ready: 5 6 7 8 1 9 0 2 3
Process 9    -> CPU: 4,    10:    5    Wait:
-----

Time: 0
CPU Process: 0 -> 3
Ready: 1 2 3 4 5 6 7 8 9
Wait:

Time: 9998
CPU Process: 5 -> 7
I/O Process: 4 -> 8
Ready: 6 7 8 1 9 0 2 3
Wait: 4

Time: 1
CPU Process: 0 -> 2
Ready: 1 2 3 4 5 6 7 8 9
Wait:

Time: 9999
CPU Process: 5 -> 6
I/O Process: 4 -> 7
Ready: 6 7 8 1 9 0 2 3
Wait: 4

Time: 2
CPU Process: 0 -> 1
Ready: 1 2 3 4 5 6 7 8 9
Wait:

Time: 10000
CPU Process: 5 -> 5
I/O Process: 4 -> 6
Ready: 6 7 8 1 9 0 2 3
Wait: 4

Time: 3
CPU Process: 1 -> 3
I/O Process: 0 -> 5
Ready: 2 3 4 5 6 7 8 9
Wait: 0

Result
Troughput: 0.178
Average Wait Time: 51.472
Average Turnaround Time: 45.871
-----
```

- Result of the simple scheduler with SJF policy and 10 maximum burst time:

Process 0	-> CPU: 1,	IO: 9	Time: 9997
Process 1	-> CPU: 5,	IO: 7	CPU Process: 5 -> 1
Process 2	-> CPU: 3,	IO: 3	I/O Process: 0 -> 2
Process 3	-> CPU: 6,	IO: 2	I/O Process: 4 -> 5
Process 4	-> CPU: 2,	IO: 10	I/O Process: 7 -> 3
Process 5	-> CPU: 2,	IO: 4	Ready: 2 8 6 3 1 9
Process 6	-> CPU: 4,	IO: 2	Wait: 0 4 7
Process 7	-> CPU: 3,	IO: 5	
Process 8	-> CPU: 5,	IO: 4	Time: 9998
Process 9	-> CPU: 9,	IO: 5	CPU Process: 2 -> 3
			I/O Process: 0 -> 1
			I/O Process: 4 -> 4
			I/O Process: 7 -> 2
			I/O Process: 5 -> 3
			Ready: 6 8 9 3 1
			Wait: 0 4 7 5
Time: 0			
CPU Process: 0 -> 1			
Ready: 4 7 5 8 1 2 6 3 9			
Wait:			
Time: 1			Time: 9999
CPU Process: 4 -> 2			CPU Process: 2 -> 2
I/O Process: 0 -> 8			I/O Process: 4 -> 3
Ready: 5 7 2 8 1 9 6 3			I/O Process: 7 -> 1
Wait: 0			I/O Process: 5 -> 2
			Ready: 0 8 6 3 1 9
			Wait: 4 7 5
Time: 2			
CPU Process: 4 -> 1			Time: 10000
I/O Process: 0 -> 7			CPU Process: 2 -> 1
Ready: 5 7 2 8 1 9 6 3			I/O Process: 4 -> 2
Wait: 0			I/O Process: 5 -> 1
			Ready: 0 8 7 3 1 9 6
			Wait: 4 5
Time: 3			
CPU Process: 5 -> 2			Result
I/O Process: 0 -> 6			Troughput: 0.463
I/O Process: 4 -> 9			Average Wait Time: 4.880
Ready: 7 8 2 3 1 9 6			Average Turnaround Time: 2.720
Wait: 0 4			

- Result of the simple scheduler with RR policy and 10 maximum burst time:

Process 0	-> CPU: 8,	IO: 8	Time: 9996
Process 1	-> CPU: 5,	IO: 4	CPU Process: 2 -> 3
Process 2	-> CPU: 6,	IO: 6	Ready: 4 0 3 7 5 9 1 6 8
Process 3	-> CPU: 1,	IO: 1	Wait:
Process 4	-> CPU: 10,	IO: 9	
Process 5	-> CPU: 5,	IO: 7	Time: 9997
Process 6	-> CPU: 7,	IO: 3	CPU Process: 2 -> 2
Process 7	-> CPU: 9,	IO: 9	Ready: 4 0 3 7 5 9 1 6 8
Process 8	-> CPU: 1,	IO: 4	Wait:
Process 9	-> CPU: 3,	IO: 6	
			Time: 9998
			CPU Process: 4 -> 10
			Ready: 0 3 7 5 9 1 6 8 2
			Wait:
Time: 0			
CPU Process: 0 -> 8			
Ready: 1 2 3 4 5 6 7 8 9			
Wait:			
Time: 1			Time: 9999
CPU Process: 0 -> 7			CPU Process: 4 -> 9
Ready: 1 2 3 4 5 6 7 8 9			Ready: 0 3 7 5 9 1 6 8 2
Wait:			Wait:
Time: 2			
CPU Process: 0 -> 6			Time: 10000
Ready: 1 2 3 4 5 6 7 8 9			CPU Process: 4 -> 8
Wait:			Ready: 0 3 7 5 9 1 6 8 2
			Wait:
Time: 3			
CPU Process: 0 -> 5			Result
Ready: 1 2 3 4 5 6 7 8 9			Troughput: 0.217
Wait:			Average Wait Time: 40.811
			Average Turnaround Time: 36.206

8 Evaluation

Process 0	-> CPU: 7,	10:	4
Process 1	-> CPU: 6,	10:	8
Process 2	-> CPU: 6,	10:	4
Process 3	-> CPU: 3,	10:	7
Process 4	-> CPU: 2,	10:	10
Process 5	-> CPU: 8,	10:	3
Process 6	-> CPU: 10,	10:	1
Process 7	-> CPU: 7,	10:	4
Process 8	-> CPU: 7,	10:	1
Process 9	-> CPU: 7,	10:	3

Figure 34 - Generated Job Processes List

Result
Throughput: 0.159
Average Wait Time: 58.331
Average Turnaround Time: 52.034

Figure 35 - Result of the Scheduling Done by FCFS Policy

Result
Throughput: 0.172
Average Wait Time: 52.879
Average Turnaround Time: 47.086

Figure 36 - Result of the Scheduling Done by RR Policy

Result
Throughput: 0.235
Average Wait Time: 9.750
Average Turnaround Time: 5.500

Figure 37 - Result of the Scheduling Done by SJF Policy




Figure 38 - Representation of the Normalized Scheduling Criteria for Different Scheduling Algorithms

Figures 35 and 36 show the graph result of execution time per number of the threads for the producer and consumer program versions 2.3 that reads the file of FeeBSD9-Orig.tar and the android.tar. The following program is implemented by applying the methods that are presented previously. As we can see from the figures, the execution time decreases as the number of threads increases. In short, the producer and consumer program version 2.3 show the ideal and faster result of execution time among the other programs.

9 Conclusion

By understanding this paper, we can understand the basic concepts of thread and multi-threaded programming. Also, we can understand the problems and the solutions that occur by applying the following concept, which is about data dependency and synchronization.

Citations

- [1] Kirvan, P. (2022, May 26). *What is multithreading?* WhatIs.com. Retrieved September 11, 2022, from <https://www.techtarget.com/whatis/definition/multithreading>
- [2] Silberschatz, A., Galvin, P. B., & Gagne, G. (2014). 2.2.1 Command Interpreters. In *Operating Systems Concepts*. essay, Wiley.
- [3] Yan, D. (2020, December 22). *Producer-consumer problem using mutex in C++*. Medium. Retrieved September 9, 2022, from <https://levelup.gitconnected.com/producer-consumer-problem-using-mutex-in-c-764865c47483>
- [4] Yoo, S. H. (n.d.). *Mobile-os-DKU-cis-MSE*. GitHub. Retrieved September 8, 2022, from <https://github.com/mobile-os-dku-cis-mse/>